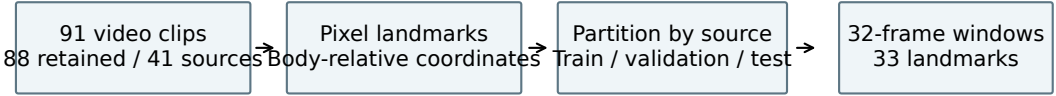
91 video clips
88 retained / 41 sources
Pixel landmarks
Body-relative coordinates
Partition by source
Train / validation / test
32-frame windows
33 landmarks
Training sources: label-free encoder + supervised probe
Test sources: one prediction per clip
Validation: choose 800 or 800+400 updates
Pool scores with equal source or equal clip weights