

Lane C improved because its evaluation definition was repaired

The S-JEPA checkpoint and all 159 saved embeddings stayed fixed.

Superseded: 5 ordinary group folds

accuracy 0.604
balanced accuracy 0.595
macro-F1 0.407

One training fold had no cerebral-palsy rows.
Macro-F1 label sets varied by fold.

same encoder
new folds

Corrected: 2 stratified group folds

accuracy 0.653 (change +0.048)
balanced accuracy 0.603 (change +0.008)
macro-F1 0.625 (change +0.218)

Every train and test fold contains all five labels.
Macro-F1 always uses one fixed label order.

Why only two folds? Parkinson's and cerebral palsy have two source videos each. Two is the largest fold count that keeps every class on both sides.

What did not improve: the encoder was trained once on all 159 rows. Classifier-level video grouping does not create an independent encoder test.