

One frozen predictor, three mask geometries

(a) training: spatial infilling



16 time patches

hide random joints (12-landmark
whitelist), full time context

(b) world-model probe: future masking



16 time patches

hide ALL joints of the last h patches:
predict the future in latent space

(c) what the frozen model achieves

spatial infilling (native task)	0.547
future, h=2 (Parkinson's, best)	0.442
future, h=2 (cerebral palsy, worst)	0.233
copy-last-patch baseline	0.88-0.95

(d) 2-step latent rollout (one normal clip)

direct h=2 cosine 0.571 -> chained-rollout cosine 0.608
(imagined patch feeds the next prediction; error did not explode in 2 steps)

verdict: the infilling predictor does NOT forecast better than a memoryless baseline - infilling != forecasting