

What is implemented, proposed, or explicitly rejected

Implemented and measured

- 12-landmark fixed target whitelist
- uniform batch-safe masking with validity gates
- 63 accepted added-normal pose sequences
- five stages with replay, VICReg, and group loss
- fingerprinted checkpoints, embeddings, and result ledger
- A1, A2, and encoder-transductive Lane C audits

Proposed next experiments

- outer source-video split before all fitting
- fold-local five-stage representation training
- source-balanced acquisition and provenance controls
- dynamics and raw-rate readout ablations
- same-pose handcrafted, embedding, and fused comparison
- uncertainty from genuinely independent outer folds

Rejected interpretations

- Lane C is an unseen-video encoder estimate
- every sample masks exactly 60 percent
- added normal labels are canonical GAVD labels
- VICReg guarantees clean clinical clusters
- binary and five-class bar heights are comparable
- a fixed Random Forest can undo encoder exposure

The archived improvement plan remains useful for hypotheses, but only the green column describes the completed checkpoint.