



The current fold-local model



At most 12 eligible joints × 16 segments = 192 target positions.
Mask 60% of the batch-minimum valid-eligible count.
Fully valid case: 115 targets / 528 tokens = 21.8%.

Both encoders receive the same prepared sequence.



$$\text{Loss} = \text{Smooth-L1} + 0.10 \text{ variance} + 0.01 \text{ covariance}$$