Emergent vs built-in: two ways to get the mirror rule

Two ways a model could satisfy the mirror rule s(Mx) = -s(x): build it in by wiring (top), or hope training teaches it (bottom).

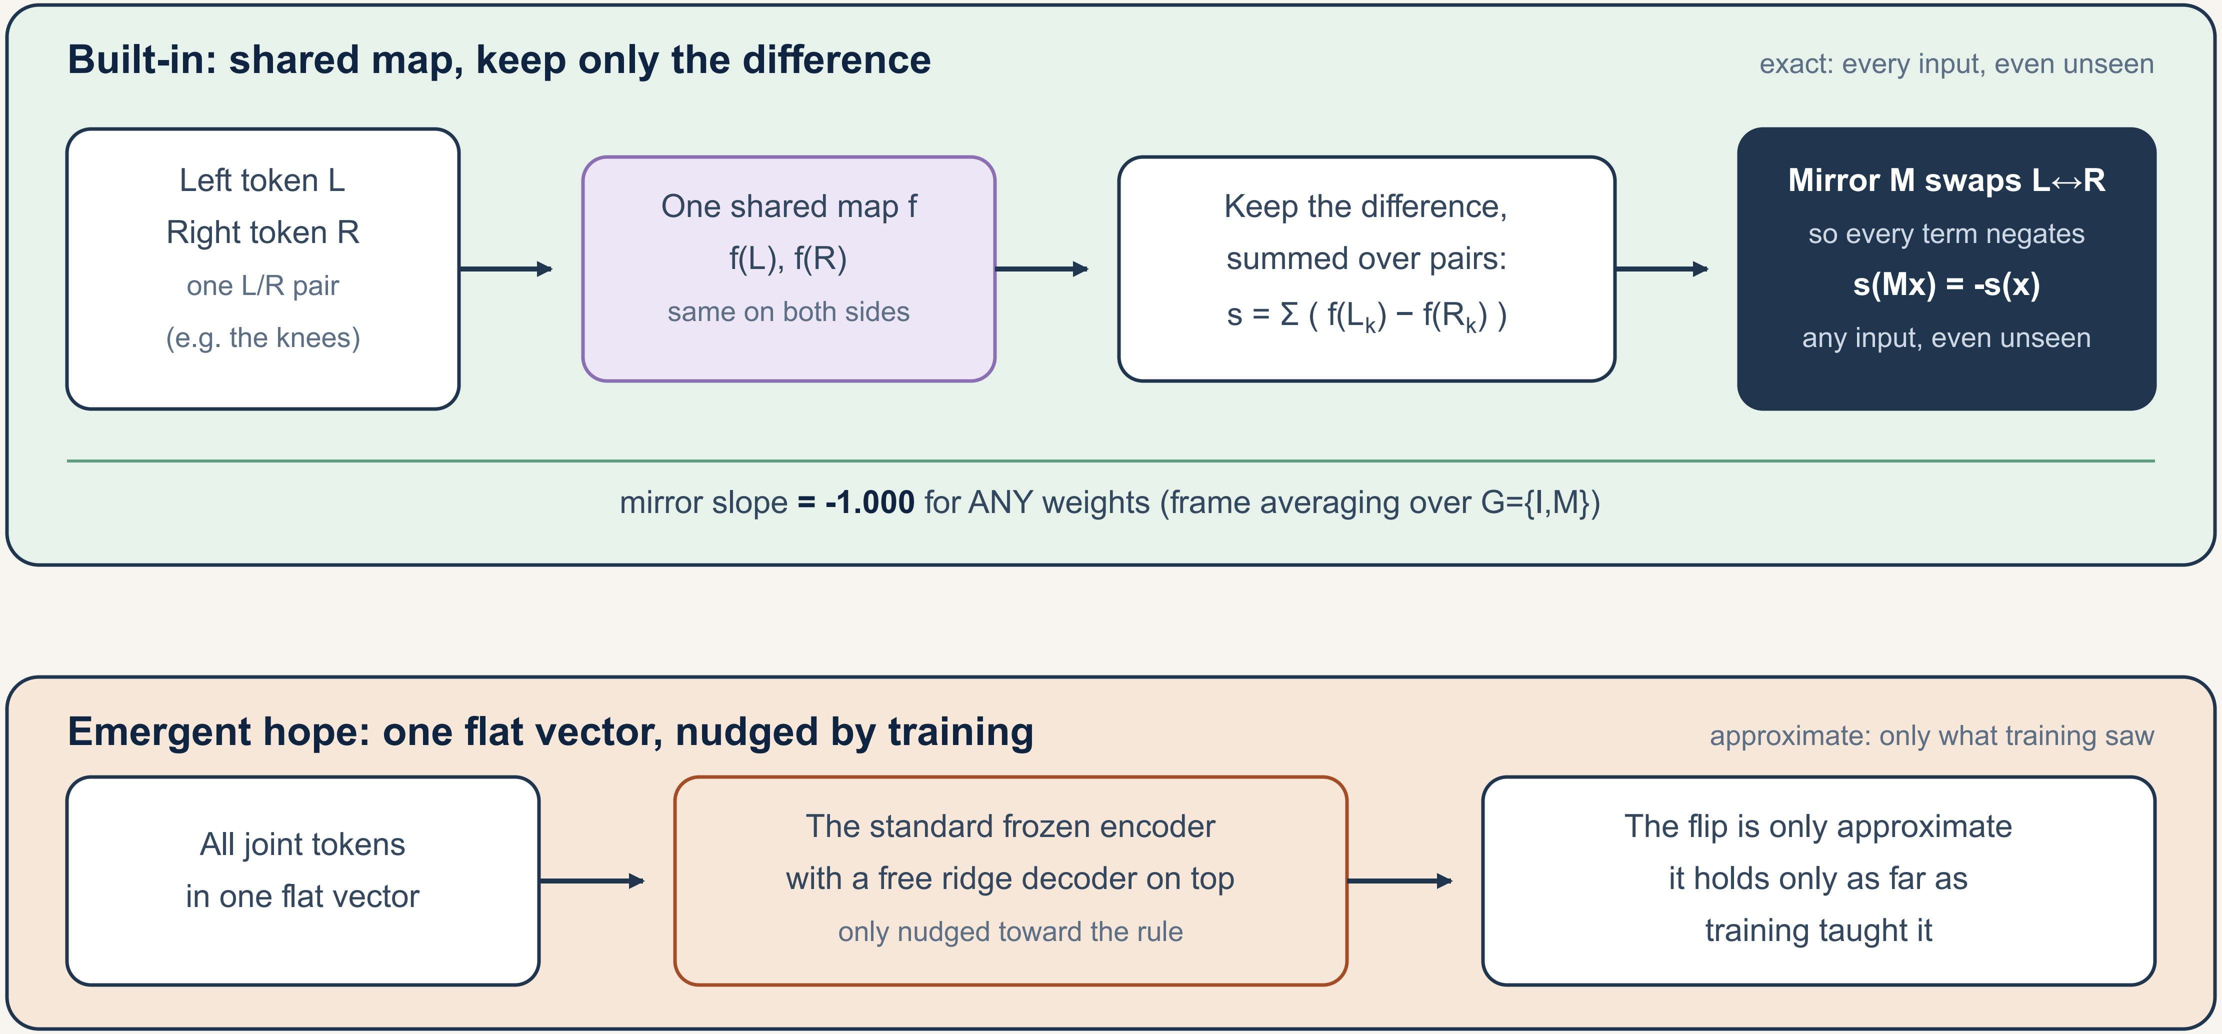

Reflection augmentation is OFF in canonical training (flip probability 0.0): this is a genuine test of the property, not something taught.

Illustrative schematic. Folder labels are dataset annotations, not diagnoses.