


SOURCE-VIDEO SPLIT | every clip and derived view keeps its source assignment



Timestamped observations

$T \times 33 \times (x, y, z, \text{visibility})$

● observed ○ missing

Schematic poses; no forecast

Movement-measure path

Observed transitions + original time

Five paired median speeds $\rightarrow$ signed y

Encoder-input path

Fill ≤ 4 samples; normalize
Resize by index: $64 \times 33 \times 3$
Original time intervals omitted
Validity kept separately

Four-position patches

16 blocks $\times$ 33 joints



Context can flank targets

Online encoder

Predictor

mask tokens

Hidden feature loss

CE

Teacher

no gradient

EMA

Full prepared clip to teacher

Two full geometric views $\rightarrow$ same online encoder $\rightarrow$ 12-joint pool $\rightarrow$ VICReg

FROZEN EVALUATION

Held-out source videos

Initial / trained encoder
paired feature summary

Ridge readout
fit on train sources

Estimate y
score against y