

The current 384-D readout is simple, but it compresses away temporal order

The right side is a proposed controlled ablation. It has not produced a saved result in this repository.

Encoded token grid

16 time patches
x 33 joints
x 96 latent values

valid tokens only



Current pooling

global mean: 96
global std: 96
12-landmark mean: 96
12-landmark std: 96

Total = 384 values



Planned dynamics ablation

per-segment profiles
left-right trajectories
raw duration and frame rate
cadence or autocorrelation
fit inside each outer fold

PROPOSED

Why the current design is useful

No trainable sequence head.
Low variance at tiny sample size.
Easy-to-audit frozen feature probe.

What it cannot retain

Order of the 16 time patches.
Absolute walking rate after resizing.
When left-right asymmetry occurs.