The current masking rule is constrained, uniform, and batch-safe

The numbers in the middle card are a teaching example. Saved run values appear in the final card.

1. Build joint-time tokens

64 normalized frames
÷ 4 frames per patch
= 16 time patches

16 x 33 joints
= 528 possible tokens

All 33 identities may provide context.
Only 12 may become targets.

2. Worked batch example

Valid eligible counts:
180, 160, 150, 175

common target count
= floor(0.60 x 150)
= 90 tokens per sample

realized fractions:
0.500, 0.562, 0.600, 0.514

3. What the run recorded

Stage 0 endpoint mean: 0.551
Stage 4 endpoint mean: 0.423

A sample with more eligible tokens gets a lower realized fraction.

Targets are drawn uniformly.
No velocity, displacement, or learned motion score is consulted.

This corrects the misleading shortcut phrase '60% of each sample.' The setting is applied once to the smallest eligible count in the batch.